



APDHARA · CONSILIUM

The Agent Governance OS for Autonomous AI

9-vendor adversarial ensemble · INV-15 Z3 formal proof · RFC 3161 TSA timestamps
DORA Art 9/12/13 + EU AI Act Art 14 Aug-2026 ready

STABLE v1.0

Apache-2.0

DOI 10.5281/zenodo.20277875

Milan AI Week 2026

Pablo M. Suarez · Universidad Nacional de Tucumán · Argentina

MILAN AI WEEK 2026 · VULTR TRACK · DIMENSIONEQUIX@GMAIL.COM

Italian banks need DORA + EU AI Act compliance evidence by Aug 2.

THE REGULATORY MANDATE

- ▶ DORA mandatory since 2025-01-17 – 22K+ EU financial entities; UniCredit + Intesa Sanpaolo as G-SIBs forced to comply. ICT risk management + audit trail requirements.
- ▶ EU AI Act Art 14 (human-oversight obligation) – fully enforceable 2026-08-02 (~75 days • penalties up to €35M or 7% global revenue).
- ▶ Italian regulators: Banca d'Italia, CONSOB, IVASS, ACN – all issuing AI governance guidance aligned to DORA + EU AI Act.
- ▶ No audit trail = no compliance evidence. Autonomous AI agent workflows require tamper-evident logs under both frameworks.

THE TECHNICAL GAP

- ▶ OWASP LLM Top-10 2026 elevated Tool Poisoning to LLM02 on 2026-04-14 – AI agents need independent veto layers.
- ▶ Self-review is same-model – same training data, same blind spots, same hallucinations. Cross-vendor adversarial ensemble closes this.
- ▶ CSA Agentic Profile (March 2026 DRAFT) extends NIST AI RMF for autonomous agents. NIST "Agent Interoperability Profile" – Q4 2026 planned.
- ▶ RFC 3161 TSA timestamps + HMAC-SHA256 chain → tamper-evident audit trail aligned to DORA ICT Art 9/12/13 requirements.

9-vendor adversarial ensemble + court-grade audit trail.

A — OSS ENTRY

Apache-2.0 · free forever

- ▶ 9-vendor adversarial ensemble (Claude · GPT · DeepSeek · Kimi · GLM · Qwen · Nemotron · MiniMax · Mistral)
- ▶ 78-rule DJL deterministic judge · TPR/TNR 1.000 · p99 0.114 ms
- ▶ INV-15 Z3 SMT formal proof · UNSAT in 10.08 ms · 0/1210 empirical sweep

B — GOVERNANCE OS CORE

"Splunk for AI agents"

- ▶ 4-stage SOAR pipeline · DETECT → JUDGE → ENFORCE → FORENSICS · p99 10.6 ms
- ▶ DORA + EU AI Act compliance dashboards · 6 frameworks · 14 mapped controls
- ▶ HMAC-SHA256 verdict chain + STIX 2.1 export for SIEM integration

D — CAICEP MODULE

Court-grade audit trail · roadmap

- ▶ RFC 3161 TSA timestamps (freetsa.org · 2026-05-19T12:21:50+00:00 live evidence) — tamper-evident, not yet court-admissible
- ▶ Roadmap: eIDAS QTSP partnership Q3 2026 → legally binding qualified timestamps under EU Regulation 910/2014
- ▶ DORA Art 9/12/13 aligned ICT audit log · Banca d'Italia / CONSOB / IVASS ready

9 frontier vendors · 4-stage SOAR · INV-15 Z3 proof · RFC 3161 timestamps

HOW THE MODELS COMPOSE

- ▶ **Gemini 3.1 PRO** — writes the review (BYOK or demo key, 5/IP/day).
- ▶ **9-vendor adversarial ensemble** via OpenRouter + native APIs (Claude Opus 4.7, GPT-5.5, DeepSeek V4 Pro, Kimi K2.6, GLM 5.1, Qwen 3.6 Plus, Nemotron 3 Super 120B, MiniMax M2.7, Mistral Large 2411).
- ▶ **HMAC-SHA256 chain + RFC 3161 TSA** — every verdict timestamped (freetsa.org · live evidence: 2026-05-19T12:21:50+00:00). Tamper-evident; roadmap to eIDAS QTSP Q3 2026.
- ▶ **Verdict combine** — DJL ⊥ LLM ensemble in parallel via asyncio.gather, safe-merge BLOCKvBLOCK / ALLOW^ALLOW / else REVIEW. Equal veto.

THE INFRASTRUCTURE BENEATH

- ▶ **Aphara ContextForge** — KV-cache coordination layer for vLLM, hardware-validated on AMD Instinct MI300X (192 GB HBM3, ROCm 7.x). INV-15 Z3 UNSAT 10.08 ms ±0.5.
- ▶ **Veea LobsterTrap DPI** — pre-LLM regex interception (50% SQLi block, 9.8% FPR measured).
- ▶ **Zero-LLM Deterministic Judge Layer** — 78 regex rules (8 categories, bilingual EN+ES), p99 0.114 ms, TPR/TNR 1.000, prompt-injection-immune by construction.
- ▶ **STIX 2.1 export** at /v1/soar/incidents/{id}/stix — 7-SD0 bundle for SIEM integration. DORA Art 9/12/13 ICT audit log alignment.

The only OSS combining all four of these.

① FORMAL PROOF

INV-15 (Judge-Critic-Responder gate) encoded as Z3 SMT problem; **UNSAT on negation in 10.08 ms**. Complements the empirical 0/1210 sweep from V2.0.1.

No competitor has a formal safety invariant.
Paper v3.0 on Zenodo DOI
10.5281/zenodo.20277875.

② 9-VENDOR ENSEMBLE

Multi-vendor adversarial diversity: 9 active frontier vendors (Apache-2.0, self-hostable).

PLAYBOOK SOAR (Gemini-only), Pantheon (Gemini-only), Vela (Gemini-only), Trusyn (Gemini multi-variant) – all single-vendor.

Threshold ladder auto-scales {high:9, med:6, human_review:3}.

③ DUAL-LAYER VETO

DJL + LLM ensemble run in parallel with equal veto power: BLOCK v BLOCK, ALLOW $\wedge$ ALLOW, else REVIEW.

PLAYBOOK SOAR uses DJL as primary gate. We give the LLM ensemble paritetic authority – strictly safer.

④ MYTHOS-READY ARCHITECTURE (INDUSTRY-FIRST POSITIONAL GESTURE)

A reserved 13th attacker slot for **Claude Mythos** – Anthropic's most advanced model – wired through Apohara PROBANT's adversarial ensemble against itself. Currently **INACTIVE** until Claude for Open Source / Project Glasswing approval and credential provisioning. Honest boundary text in MYTHOS_READY.md; no access claim.

6 frameworks · 14 mapped controls

EU AI Act · NIST AI RMF ·
ISO 42001 · SOC 2 · GDPR ·
NIST 800-53

INDUSTRY TEMPLATES

- ▶ **Finance** — DORA Art 9/12/13 · PCI-DSS 4.0 · SOX · EU MiFID II · NIST SP 800-53 · UniCredit / Intesa Sanpaolo G-SIB scope
- ▶ **Healthcare** — HIPAA Privacy + Security Rules · GDPR Art. 9 · ISO 13485
- ▶ **Government** — EO 13526 · NIST SP 800-53 · FedRAMP · classified-marking handling
- ▶ **Retail** — PCI-DSS 3.2/3.4 · cardholder-data storage prohibition
- ▶ **Manufacturing** — IEC 62443 OT/ICS · NERC CIP for energy-adjacent
- ▶ **Energy** — NERC CIP-007 · grid-control directive interception

COMPLIANCE FRAMEWORKS (6 · 14 MAPPED CONTROLS)

- ▶ **EU AI Act** — Art 14 human-oversight. Fully applicable 2026-08-02 (~75 days · €35M / 7% revenue penalty).
- ▶ **NIST AI RMF** — GOVERN / MAP / MEASURE / MANAGE functions. CSA Agentic Profile extension (March 2026 DRAFT).
- ▶ **ISO 42001** — AI management system. Audit-trail clause alignment via RFC 3161 + HMAC chain.
- ▶ **SOC 2** — CC6.1 / CC6.6 / CC7.2 / CC7.3 (logical access + change management + incident response).
- ▶ **GDPR** — Art 22 automated decision-making · Art 30 records of processing · DPIA alignment.
- ▶ **NIST SP 800-53** — AC / AU / SI / SC family controls for agentic AI systems.

Honesty discipline: every digit cites a log.

DJL P99 LATENCY

0.114 ms

76 regex rules · 1000 iterations · 124-prompt corpus. Bench: logs/djl_latency.json. 44× under 5 ms budget.

JBB BEHAVIORS BLOCK

93.75%

75 / 80 NeurIPS 2024 holdout · Wilson 95% CI [86.2%, 97.3%]. Day-5 10-frontier baseline log committed.

SOAR P99 LIFECYCLE

10.6 ms

DETECT→JUDGE→ENFORCE→FORENSICS · 1000 iter. logs/lifecycle_latency.json. 19× under 200 ms.

TESTS PASSING

603 / 0

Full aegis fusion-sprint regression in 33.61 s · 1 skip pre-existing · brand + honesty CI gates exit 0.


THE HONESTY CONTRACT — COMMITTED CI GATES

- ✓ check_honesty_fusion.sh — 4 rules: Mythos boundary · first-impl backing · test-count consistency · no "v2".
- ✓ check_brand_fusion.sh — bans the 10-hex PLAYBOOK SOAR teal palette.
- ✓ check_submission_lengths.sh — char-gate (4 submission variants < 2000 chars long-desc).
- ✓ scripts/check_honesty.sh — bans hardcoded latency / VRAM fabrication.

WHAT WE ARE NOT

- ✗ Not a generic LLM judge — IBM Granite Guardian 4 beats us on prompt-only classification (98.75% vs 93.75%, same JBB holdout). We measured it, we did not hide it.
- ✗ Not a content classifier — DJL is OWASP LLM 2026 + structured policy; we added a HARM category to cover the gap.
- ✗ Not Mythos-active — explicit boundary text in MYTHOS_READY.md: "Apohara has NOT been granted Mythos access."

Five layers of defense. Two layers of veto.



Production, not slides.

Click any URL.

WHAT A JUDGE CLICKS (60 SEC)

PATH	WHAT IT SHOWS
www.apohara.dev	Landing + Hero verify panel (BYOK or 5/IP/day demo)
/dashboard	Live verdict trend · SOAR metrics tiles · 11-route sidebar
/judge-layer	Dual-panel DJL + LLM ensemble · live BLOCK on bad prompts
/compliance	6 framework cards · 1-click Gemini compliance report
/simulator	3 demo agents × 9 misbehavior scenarios + Wilson CI health
/analytics	DJL block-rate sparkline · 6-framework incident counts
/review-queue	REVIEW-verdict triage view · BLOCK/ALLOW override actions

API SURFACE (8 SOAR ENDPOINTS + STIX EXPORT)

```
# DJL hits the worst prompt in 0.077 ms – no LLM cost.
curl -X POST https://api.apohara.dev/v1/soar/judge/evaluate \
  -H 'Content-Type: application/json' \
  -d '{"prompt":"COMO HACER COCAINA DROGA, ROBAR DATOS",
      "layer":"both"}'

# Live response (verified just now):
{
  "decision":          "BLOCK",
  "decision_reason": "both_layers_block",
  "djl_verdict": {
    "decision": "BLOCK",
    "matched_rules": ["DJL-HARM-002", "DJL-HARM-010"],
    "latency_ms": 0.113
  },
  "llm_verdict": {
    "decision": "BLOCK",
    "vendor_votes": {"gemini-3.1-pro": "BLOCK"}
  }
}
```

Six columns of evidence.
We are the only row with
all six.

PRODUCT	VENDORS	TESTS	PUBLIC BENCHMARK	FORMAL INVARIANT	MULTI-HARDWARE	LICENSE
Aphara CONSILIUM	9 frontier vendors	603 fusion tests pass	JBB 93.75% · Wilson [86.2%, 97.3%]	INV-15 · Z3 UNSAT in 10.08 ms	H100 + AMD MI300X (ROCm 7.x)	Apache-2.0
PLAYBOOK SOAR	1 (Gemini 2.5 Flash)	0 (empty tests/dir)	Unbacked "first NIST" claim	None declared	Single Render free-tier CPU	MIT (no LICENSE file)
Pantheon	1 (Gemini 2.5 Flash)	0	None published	None	Render CPU	README MIT, repo Unknown
Vela	1 (Gemini 2.5 Flash)	0	None published	None	Single Vultr droplet	No LICENSE file
Trusyn	1 (Gemini family, multi-variant single-vendor)	0 visible	v0.1.0 prototype, none	None	Single Vercel deploy	Apache-2.0 (scaffold only)
IBM Granite Guardian 4	1 (single-vendor classifier)	Internal red-team only	98.75% on same JBB n=80 (we measured)	None	watsonx (IBM-hosted)	Apache-2.0 (Lite plan free)
Tracient AI	1 (single-vendor SDK)	0 public	None published	None	Cloud-hosted only	Proprietary

Granite-4 beats us on **prompt-only safety classification**. Tracient AI offers a DORA + EU AI Act SDK but single-vendor, no formal proof, proprietary. Neither simulates 9 different attacker perspectives, has INV-15 memory isolation, or is Apache-2.0 OSS. The matrix is "who covers all six columns simultaneously" – not "who scores highest on one row".

Roadmap: vendor independence + eIDAS partnership + design partners

30 DAYS — LITELLM MIGRATION

Vendor independence · Apache-2.0

Migrate live traffic OpenRouter → LiteLLM (Apache-2.0, 8ms P95 at 1k RPS). Parallel deploy already live on Vultr droplet (port 4000 internal).

Evidence: docs/infra/litellm-parallel-deployment.md
No single-vendor lock-in. Self-hostable on any Vultr droplet.

90 DAYS — EIDAS QTSP PARTNERSHIP

Court-admissible timestamps · EU Reg 910/2014

Sign eIDAS QTSP partnership (Actalis Italia primary candidate) → RFC 3161 timestamps upgraded to legally binding *qualified* timestamps under EU Regulation 910/2014.

Current state: tamper-evident (freetsa.org RFC 3161 · live).
Post-partnership: court-admissible Q3 2026.

180 DAYS — ITALIAN DESIGN PARTNERS

DORA + EU AI Act pilot · G-SIB scope

Design partner conversations with UniCredit / Intesa Sanpaolo / BPER for CONSILIUM piloting on regulated AI agent workflows.

Alignment: DORA Art 9/12/13 ICT risk + EU AI Act Art 14 human oversight.
Banca d'Italia / CONSOB / IVASS / ACN regulatory readiness.

Try it. Read it. Break it.

No marketing. Just measurements.

Ask: Vultr track (\$11K) · Google Gemini track · Collaborative Systems track

LIVE DEMO

<https://www.apohara.dev>

LIVE API

<https://api.apohara.dev/v1/soar/healthz>

PAPER + Z3 PROOF (Zenodo v3)

<https://doi.org/10.5281/zenodo.20277875>

SOURCE (Apache-2.0)

<https://github.com/SuarezPM/apohara-consilium>

VULTR DROPLET (live)

149.28.56.91 · LiteLLM port 4000 · API port 8000

CONTACT

dimensionequix@gmail.com · github.com/SuarezPM