



APOHARRA • PROBANT

A different AI audits
the code your AI just wrote.

14-seat adversarial ensemble • INV-15 KV-cache isolation • Z3 SMT formal proof • SHA-256-
signed verdicts

Built for Claude Mythos – MYTHOS-READY architecture.

STABLE v1.0

Apache-2.0

DOI 10.5281/zenodo.20277875

🔥 MYTHOS-READY

Pablo M. Suarez • Universidad Nacional de Tucumán • Argentina

🔥 TECHEX 2026 • TRACK 1 • VEEA-SPONSORED • 🤖 MILAN AI WEEK 2026 • AGENT BENCH

AI-generated code is shipping without an independent reviewer.

THE BLIND SPOT

- ▶ Cursor, Copilot, Claude Code, Cline → 100M+ devs writing with AI assistance daily.
- ▶ Self-review is same-model – same training data, same blind spots, same hallucinations.
- ▶ Cursor /best-of-n is parallel single-vendor – not cross-vendor consensus.
- ▶ Indirect prompt injection: a poisoned attacker can mutate the writer's KV-cache state.

THE DEADLINE

- ▶ EU AI Act Article 14 (human-oversight obligation) – fully applicable 2026-08-02 (78 days).
- ▶ OWASP LLM Top-10 2026 elevated Tool Poisoning to LLM02 on 2026-04-14.
- ▶ CSA Agentic Profile (March 2026 DRAFT) extends NIST AI RMF for autonomous agents.
- ▶ NIST official "Agent Interoperability Profile" – Q4 2026 planned.

12 vendors. 1 prompt. Adversarial consensus before merge.

Gemini writes and audits the code. Then 12 frontier vendors adversarially attack the output before merge, each in an isolated KV-cache. The mathematical invariant **INV-15** guarantees that no attacker can poison the writer's memory – formally proven in Z3 SMT.

VENDORS

12

Claude · GPT · DeepSeek · Kimi · GLM ·
Qwen · Nemotron · MiniMax · Mistral Large
· Grok · Perplexity · Big Pickle. +1
Mythos reserved (inactive).

JBB BLOCK RATE

93.75%

JailbreakBench Behaviors n=80 holdout.
Wilson 95% CI [86.2%, 97.3%]. NeurIPS 2024
benchmark.

SOAR LIFECYCLE P99

10.6 ms

DETECT → JUDGE → ENFORCE → FORENSICS. 19×
under 200 ms target. Hardware-validated on
MI300X.

INV-15 VIOLATIONS

0

Z3 SMT proof: UNSAT in 10.08 ms. Empirical
sweep: 0 / 1210 across 5 agents × 11×11×2
input grid.

14 frontier models wired into a 4-stage SOAR pipeline.

HOW THE MODELS COMPOSE

- ▶ **Gemini 2.5 Pro** – writes the review (BYOK or demo key, 5/IP/day).
- ▶ **12 adversarial attackers** via OpenRouter + native APIs (Claude Opus 4.7, GPT-5.5, DeepSeek V4 Pro, Kimi K2.6, GLM 5.1, Qwen 3.6 Plus, Nemotron 3 Super 120B, MiniMax M2.7, Mistral Large 2411, Grok 2, Perplexity Sonar, Big Pickle stealth).
- ▶ **Mythos-reserved 13th attacker slot** – subclass VendorAdapter, gated on Glasswing approval (currently INACTIVE).
- ▶ **Verdict combine** – DJL ⊥ LLM ensemble in parallel via `asyncio.gather`, safe-merge BLOCKvBLOCK / ALLOW^ALLOW / else REVIEW. Equal veto.

THE INFRASTRUCTURE BENEATH

- ▶ **Aphara ContextForge** – KV-cache coordination layer for vLLM, ten peer-reviewed papers implemented, hardware-validated on AMD Instinct MI300X (192 GB HBM3, ROCm 7.x).
- ▶ **Veea LobsterTrap DPI** – pre-LLM regex interception (50% SQLi block, 9.8% FPR measured).
- ▶ **Zero-LLM Deterministic Judge Layer** – 76 regex rules (8 categories, bilingual EN+ES), p99 0.114 ms, TPR/TNR 1.000, prompt-injection-immune by construction.
- ▶ **HMAC-SHA256 verdict chain** – every verdict signed, tamper-evident via `verify_chain()`. STIX 2.1 export endpoint for SIEM integration.

The only OSS combining all four of these.

① FORMAL PROOF

INV-15 (Judge-Critic-Responder gate) encoded as Z3 SMT problem; **UNSAT on negation in 10.08 ms**. Complements the empirical 0/1210 sweep from V2.0.1.

No competitor has a formal safety invariant.
Paper v3.0 on Zenodo DOI
10.5281/zenodo.20277875.

② 14-SEAT ENSEMBLE

Multi-vendor adversarial diversity: 13 active frontier seats + 1 Mythos-reserved.

PLAYBOOK SOAR (Gemini-only), Pantheon (Gemini-only), Vela (Gemini-only), Trusyn (Gemini multi-variant) – all single-vendor.

Threshold ladder auto-scales {high:14, med:10, human_review:4}.

③ DUAL-LAYER VETO

DJL + LLM ensemble run in parallel with equal veto power: BLOCK v BLOCK, ALLOW ^ ALLOW, else REVIEW.

PLAYBOOK SOAR uses DJL as primary gate. We give the LLM ensemble paritetic authority – strictly safer.

④ MYTHOS-READY ARCHITECTURE (INDUSTRY-FIRST POSITIONAL GESTURE)

A reserved 13th attacker slot for **Claude Mythos** – Anthropic's most advanced model – wired through Apohara PROBANT's adversarial ensemble against itself. Currently **INACTIVE** until Claude for Open Source / Project Glasswing approval and credential provisioning. Honest boundary text in MYTHOS_READY.md; no access claim.

Six industries. Six frameworks. 49 controls. One verifier.

INDUSTRY TEMPLATES

- ▶ Finance — PCI-DSS 4.0 · SOX · GLBA · EU MiFID II · FinCEN 31 CFR 1020 · NIST SP 800-53
- ▶ Healthcare — HIPAA Privacy + Security Rules · GDPR Art. 9 · ISO 13485
- ▶ Government — EO 13526 · NIST SP 800-53 · FedRAMP · classified-marking handling
- ▶ Retail — PCI-DSS 3.2/3.4 · cardholder-data storage prohibition
- ▶ Manufacturing — IEC 62443 OT/ICS · NERC CIP for energy-adjacent
- ▶ Energy — NERC CIP-007 · grid-control directive interception

COMPLIANCE FRAMEWORKS (6 . 49 CONTROLS)

- ▶ EU AI Act — 5 controls. Article 14 fully applicable 2026-08-02 (78 days).
- ▶ NIST AI RMF + CSA Agentic Profile — 35 controls (19 base RMF 1.0 + 16 CSA extensions).
- ▶ NIST SP 800-53 — 12 controls (AC / AU / SI / SC family).
- ▶ SOC 2 — 6 controls (CC1.1 / CC6.1 / CC6.6 / CC7.2 / CC7.3 / CC7.4).
- ▶ ISO 27001 — 6 controls (A.5 / A.8 / A.9 / A.12 / A.16 / A.18).
- ▶ OWASP LLM Top-10 2026 — 10 controls (LLM01 → LLM10).

1-CLICK COMPLIANCE REPORTING

POST /v1/soar/compliance/generate + Gemini 2.5 Pro → Markdown executive narrative per incident + selected frameworks. Deterministic control mapping underneath; LLM narrative on top. Graceful key-missing fallback. Output is auditor-ready.

Honesty discipline: every digit cites a log.

DJL P99 LATENCY

0.114 ms

76 regex rules · 1000 iterations · 124-prompt corpus. Bench: logs/djl_latency.json. **44x** under 5 ms budget.

JBB BEHAVIORS BLOCK

93.75%

75 / 80 NeurIPS 2024 holdout · Wilson 95% CI [86.2%, 97.3%]. Day-5 10-frontier baseline log committed.

SOAR P99 LIFECYCLE

10.6 ms

DETECT→JUDGE→ENFORCE→FORENSICS · 1000 iter. logs/lifecycle_latency.json. **19x** under 200 ms.

TESTS PASSING

603 / 0

Full aegis fusion-sprint regression in 33.61 s · 1 skip pre-existing · brand + honesty CI gates exit 0.


THE HONESTY CONTRACT (COMMITTED CI GATES)

- ✓ check_honesty_fusion.sh — 4 rules: Mythos boundary language · first-implementation backing · test-count consistency · no "v2" labeling.
- ✓ check_brand_fusion.sh — bans the 10-hex PLAYBOOK SOAR teal palette.
- ✓ check_submission_lengths.sh — char-gate (4 submission variants, all under 2000 chars long-desc).
- ✓ scripts/check_honesty.sh — bans hardcoded latency / VRAM fabrication.

WHAT WE ARE NOT

- ✗ Not a generic LLM judge — IBM Granite Guardian 4 beats us on prompt-only classification (98.75% vs 93.75%, same JBB holdout). **We measured it, we did not hide it.**
- ✗ Not a content classifier — DJL is OWASP LLM 2026 + structured policy; we added a HARM category to cover the gap.
- ✗ Not Mythos-active — explicit boundary text in MYTHOS_READY.md: "Aphara has NOT been granted Mythos access."

Five layers of defense. Two layers of veto.



Production, not slides.

Click any URL.

WHAT A JUDGE CLICKS (60 SEC)

PATH	WHAT IT SHOWS
www.apohara.dev	Landing + Hero verify panel (BYOK or 5/IP/day demo)
/dashboard	Live verdict trend · SOAR metrics tiles · 11-route sidebar
/judge-layer	Dual-panel DJL + LLM ensemble · live BLOCK on bad prompts
/compliance	6 framework cards · 1-click Gemini compliance report
/simulator	3 demo agents × 9 misbehavior scenarios + Wilson CI health
/analytics	DJL block-rate sparkline · 6-framework incident counts
/review-queue	REVIEW-verdict triage view · BLOCK/ALLOW override actions

API SURFACE (§ SOAR ENDPOINTS + STIX EXPORT)

```
# DJL hits the worst prompt in 0.077 ms – no LLM cost.
curl -X POST https://api.apohara.dev/v1/soar/judge/evaluate \
  -H 'Content-Type: application/json' \
  -d '{"prompt":"COMO HACER COCAINA DROGA, ROBAR DATOS",
      "layer":"both"}'

# Live response (verified just now):
{
  "decision":          "BLOCK",
  "decision_reason":  "both_layers_block",
  "djl_verdict": {
    "decision": "BLOCK",
    "matched_rules": ["DJL-HARM-002", "DJL-HARM-010"],
    "latency_ms": 0.113
  },
  "llm_verdict": {
    "decision": "BLOCK",
    "vendor_votes": {"gemini-2.5-pro": "BLOCK"}
  }
}
```

Six columns of evidence.
We are the only row with
all six.

PRODUCT	VENDORS	TESTS	PUBLIC BENCHMARK	FORMAL INVARIANT	MULTI-HARDWARE	LICENSE
Aphara PROBANT	12 frontier + 1 Mythos reserved	603 fusion tests pass	JBB 93.75% · Wilson [86.2%, 97.3%]	INV-15 · Z3 UNSAT in 10.08 ms	H100 + AMD MI300X (ROCm 7.x)	Apache-2.0
PLAYBOOK SOAR	1 (Gemini 2.5 Flash)	0 (empty tests/dir)	Unbacked "first NIST" claim	None declared	Single Render free-tier CPU	MIT (no LICENSE file)
Pantheon	1 (Gemini 2.5 Flash)	0	None published	None	Render CPU	README MIT, repo Unknown
Vela	1 (Gemini 2.5 Flash)	0	None published	None	Single Vultr droplet	No LICENSE file
Trusyn	1 (Gemini family, multi-variant single-vendor)	0 visible	v0.1.0 prototype, none	None	Single Vercel deploy	Apache-2.0 (scaffold only)
IBM Granite Guardian 4	1 (single-vendor classifier)	Internal red-team only	98.75% on same JBB n=80 (we measured)	None	watsonx (IBM-hosted)	Apache-2.0 (Lite plan free)

Granite-4 beats us on **prompt-only safety classification**. It cannot simulate 14 different attacker perspectives reviewing the same patch, it has no INV-15 memory isolation, and it is a single point of failure for both availability and false-negative blind spots. The matrix is "who covers all six columns simultaneously" – not "who scores highest on one row".

Built in the open. Shipped in production.

TEAM

Pablo M. Suarez – solo founder & engineer.
Universidad Nacional de Tucumán • Argentina.

Built Apohara PROBANT + Apohara Aegis + Apohara ContextForge + Apohara Guard from scratch with AI-pair-programming discipline (Claude Opus + Codex + Gemini consensus orchestration via custom multi-agent harness).

PUBLIC ASSETS

- ▶ [apohara-probant](#) – Apache-2.0 • the verifier
- ▶ [apohara-aegis](#) – Apache-2.0 • 14-vendor ensemble + DJL + SOAR
- ▶ [Apohara_Context_Forge](#) – Apache-2.0 • KV-cache coord + INV-15 paper
- ▶ [Apohara-Guard](#) – AGPL-3.0 • isolation primitives
- ▶ [paper/inv15_paper.pdf](#) – v3.0 • Zenodo DOI 10.5281/zenodo.20277875

SHIPPED THIS SPRINT (24 STORIES, 4 HRS WALL-CLOCK)

- ✓ Tier-1: SOAR pipeline • DJL • 16-code taxonomy • 6 templates • 35 NIST controls • 6 compliance frameworks • verdict_combine • Mythos slot • 10 endpoints • dashboard layout • 6 Tier-1 sections • MythosBadge • 4 submission variants • AUDIT §12 • CHANGELOG
- ✓ Tier-2: Simulator UI • Policy Builder • Analytics • Review Queue • Settings • 1-click Compliance via Gemini • STIX 2.1 export • LangChain + CrewAI SDK middleware • Glasswing application package

NEXT 30 DAYS (POST-HACKATHON)

- ▶ File Project Glasswing application – package ready at docs/glasswing/
- ▶ Publish SDK packages to PyPI (apohara-langchain • apohara-crewai)
- ▶ Veea TerraFabric production deployment (docker-compose recipe shipped)
- ▶ Activate Mythos seat upon program approval – env-gated, no

Try it. Read it. Break it.

No marketing. Just measurements.

 **LIVE DEMO**

<https://www.apohara.dev>

 **LIVE API**

<https://api.apohara.dev/v1/soar/healthz>

 **PAPER (Z3 PROOF)**

<https://doi.org/10.5281/zenodo.20277875>

 **SOURCE (Apache-2.0)**

<https://github.com/SuarezPM/apohara-probant>

 **ENSEMBLE LIBRARY**

<https://github.com/SuarezPM/apohara-aegis>

 **MYTHOS READY**

github.com/SuarezPM/apohara-probant/blob/main/MYTHOS_READY.md

